

Data Mining and Machine Learning

A Master Course

Master program on Data Science

Tran Trong Hieu¹

¹VNU - University of Engineering and Technology

Chapter 24: Logistic Regression

Binary Logistic Regression

Given a set of d predictor or independent variables $X_1, X_2, \dots, X_d$, and a *binary* or *Bernoulli* response variable Y that takes on only two values, namely, 0 and 1.

Since there are only two outcomes for Y , its PMF for $\tilde{\mathbf{X}} = \tilde{\mathbf{x}}$ is:

$$P(Y = 1 | \tilde{\mathbf{X}} = \tilde{\mathbf{x}}) = \pi(\tilde{\mathbf{x}}) \qquad P(Y = 0 | \tilde{\mathbf{X}} = \tilde{\mathbf{x}}) = 1 - \pi(\tilde{\mathbf{x}})$$

where $\pi(\tilde{\mathbf{x}})$ denotes the probability of $Y = 1$ given $\tilde{\mathbf{X}} = \tilde{\mathbf{x}}$.


Instead of directly predicting the response value, the goal is to learn the probability, $P(Y = 1 | \tilde{\mathbf{X}} = \tilde{\mathbf{x}})$, which is also the expected value of Y given $\tilde{\mathbf{X}} = \tilde{\mathbf{x}}$.

Since $P(Y = 1 | \tilde{\mathbf{X}} = \tilde{\mathbf{x}})$ is a probability, it is not appropriate to directly use the linear regression model.

Logistic Function

Logistic regression comes from the *logistic* function (aka *sigmoid* function) that “squashes” the output to be between 0 and 1 for any scalar input z .

$$\theta(z) = \frac{1}{1 + \exp\{-z\}} = \frac{\exp\{z\}}{1 + \exp\{z\}}$$



Logistic Function

Example

Consider what happens when z is $-\infty$, $+\infty$ and 0:

$$\theta(-\infty) = \frac{1}{1 + \exp\{\infty\}} = \frac{1}{\infty} = 0$$

$$\theta(+\infty) = \frac{1}{1 + \exp\{-\infty\}} = \frac{1}{1} = 1$$

$$\theta(0) = \frac{1}{1 + \exp\{0\}} = \frac{1}{2} = 0.5$$

$z = 0$ acts as a threshold, since, for $z > 0$, $\theta(z) > 0.5$, and, for $z < 0$, $\theta(z) < 0.5$.

Interpreting $\theta(z)$ as a probability, the larger the z value, the higher the probability.

Another interesting property of the logistic function is that

$$1 - \theta(z) = 1 - \frac{\exp\{z\}}{1 + \exp\{z\}} = \frac{1 + \exp\{z\} - \exp\{z\}}{1 + \exp\{z\}} = \frac{1}{1 + \exp\{z\}} = \theta(-z)$$

Binary Logistic Regression

We define the logistic regression model as follows:

$$P(Y = 1 | \tilde{\mathbf{X}} = \tilde{\mathbf{x}}) = \pi(\tilde{\mathbf{x}}) = \theta(f(\tilde{\mathbf{x}})) = \theta(\tilde{\omega}^T \tilde{\mathbf{x}}) = \frac{\exp\{\tilde{\omega}^T \tilde{\mathbf{x}}\}}{1 + \exp\{\tilde{\omega}^T \tilde{\mathbf{x}}\}}$$

The probability that $Y = 1$ is the output of the logistic function for the input $\tilde{\omega}^T \tilde{\mathbf{x}}$. The probability for $Y = 0$ is given as

$$P(Y = 0 | \tilde{\mathbf{X}} = \tilde{\mathbf{x}}) = 1 - P(Y = 1 | \tilde{\mathbf{X}} = \tilde{\mathbf{x}}) = \theta(-\tilde{\omega}^T \tilde{\mathbf{x}}) = \frac{1}{1 + \exp\{\tilde{\omega}^T \tilde{\mathbf{x}}\}}$$

that is, $1 - \theta(z) = \theta(-z)$ for $z = \tilde{\omega}^T \tilde{\mathbf{x}}$.

Combining these two cases the full logistic regression model is given as

$$P(Y | \tilde{\mathbf{X}} = \tilde{\mathbf{x}}) = \theta(\tilde{\omega}^T \tilde{\mathbf{x}})^Y \cdot \theta(-\tilde{\omega}^T \tilde{\mathbf{x}})^{1-Y}$$

Since Y is a Bernoulli binary random variable, $P(Y | \tilde{\mathbf{X}} = \tilde{\mathbf{x}}) = \theta(\tilde{\omega}^T \tilde{\mathbf{x}})$ when $Y = 1$ and $P(Y | \tilde{\mathbf{X}} = \tilde{\mathbf{x}}) = \theta(-\tilde{\omega}^T \tilde{\mathbf{x}})$ when $Y = 0$.

Log-Odds Ratio

Define the *odds ratio* for the occurrence of $Y = 1$ as follows:

$$\begin{aligned}\text{odds}(Y = 1|\tilde{\mathbf{X}} = \tilde{\mathbf{x}}) &= \frac{P(Y = 1|\tilde{\mathbf{X}} = \tilde{\mathbf{x}})}{P(Y = 0|\tilde{\mathbf{X}} = \tilde{\mathbf{x}})} = \frac{\theta(\tilde{\omega}^T \tilde{\mathbf{x}})}{\theta(-\tilde{\omega}^T \tilde{\mathbf{x}})} \\ &= \frac{\exp\{\tilde{\omega}^T \tilde{\mathbf{x}}\}}{1 + \exp\{\tilde{\omega}^T \tilde{\mathbf{x}}\}} \cdot (1 + \exp\{\tilde{\omega}^T \tilde{\mathbf{x}}\}) \\ &= \boxed{\exp\{\tilde{\omega}^T \tilde{\mathbf{x}}\}}\end{aligned}$$

The logarithm of the odds ratio, called the *log-odds ratio*, is therefore given as:

$$\begin{aligned}\ln(\text{odds}(Y = 1|\tilde{\mathbf{X}} = \tilde{\mathbf{x}})) &= \ln\left(\frac{P(Y = 1|\tilde{\mathbf{X}} = \tilde{\mathbf{x}})}{1 - P(Y = 1|\tilde{\mathbf{X}} = \tilde{\mathbf{x}})}\right) = \ln(\exp\{\tilde{\omega}^T \tilde{\mathbf{x}}\}) = \tilde{\omega}^T \tilde{\mathbf{x}} \\ &= \omega_0 \cdot x_0 + \omega_1 \cdot x_1 + \cdots + \omega_d \cdot x_d\end{aligned}$$

The log-odds ratio function is also called the *logit* function, defined as

$$\text{logit}(z) = \ln\left(\frac{z}{1 - z}\right)$$

It is the inverse of the logistic function.

Log-Odds Ratio

We can see that

$$\ln(\text{odds}(Y = 1 | \tilde{\mathbf{X}} = \tilde{\mathbf{x}})) = \text{logit}(P(Y = 1 | \tilde{\mathbf{X}} = \tilde{\mathbf{x}}))$$

The logistic regression model is therefore based on the assumption that the log-odds ratio for $Y = 1$ given $\tilde{\mathbf{X}} = \tilde{\mathbf{x}}$ is a linear function (or a weighted sum) of the independent attributes.

Consider the effect of attribute X_i by fixing the values for all other attributes:

$$\ln(\text{odds}(Y = 1 | \tilde{\mathbf{X}} = \tilde{\mathbf{x}})) = \omega_i \cdot x_i + C$$

$$\implies \text{odds}(Y = 1 | \tilde{\mathbf{X}} = \tilde{\mathbf{x}}) = \exp\{\omega_i \cdot x_i + C\} = \exp\{\omega_i \cdot x_i\} \cdot \exp\{C\} \propto \exp\{\omega_i \cdot x_i\}$$

where C is a constant comprising the fixed attributes.

ω_i can be interpreted as the change in the log-odds ratio for $Y = 1$ for a unit change in X_i , or, equivalently, the odds ratio for $Y = 1$ increases exponentially per unit change in X_i .

Maximum Likelihood Estimation

We will use the maximum likelihood approach to learn the weight vector $\tilde{\mathbf{w}}$.

Likelihood is defined as the probability of the observed data given $\tilde{\mathbf{w}}$.

$$L(\tilde{\mathbf{w}}) = P(Y|\tilde{\mathbf{w}}) = \prod_{i=1}^n P(y_i | \tilde{\mathbf{x}}_i) = \prod_{i=1}^n \theta(\tilde{\mathbf{w}}^T \tilde{\mathbf{x}}_i)^{y_i} \cdot \theta(-\tilde{\mathbf{w}}^T \tilde{\mathbf{x}}_i)^{1-y_i}$$

Instead of maximizing the likelihood, we can maximize the *log-likelihood*, to convert the product into a summation:

$$\ln(L(\tilde{\mathbf{w}})) = \sum_{i=1}^n y_i \cdot \ln(\theta(\tilde{\mathbf{w}}^T \tilde{\mathbf{x}}_i)) + (1 - y_i) \cdot \ln(\theta(-\tilde{\mathbf{w}}^T \tilde{\mathbf{x}}_i))$$

Maximum Likelihood Estimation

The negative of the log-likelihood can also be considered as an error function, the *cross-entropy error function*:

$$E(\tilde{\mathbf{w}}) = -\ln(L(\tilde{\mathbf{w}})) = \sum_{i=1}^n y_i \cdot \ln\left(\frac{1}{\theta(\tilde{\mathbf{w}}^T \tilde{\mathbf{x}}_i)}\right) + (1 - y_i) \cdot \ln\left(\frac{1}{1 - \theta(\tilde{\mathbf{w}}^T \tilde{\mathbf{x}}_i)}\right)$$

The task of maximizing the log-likelihood is therefore equivalent to minimizing the cross-entropy error.

Maximum Likelihood Estimation

To obtain the optimal weight vector $\tilde{\mathbf{w}}$, we would differentiate the log-likelihood function with respect to $\tilde{\mathbf{w}}$, set the result to 0, and then solve for $\tilde{\mathbf{w}}$.

However, for the log-likelihood formulation presented, there is no closed form solution to compute the weight vector $\tilde{\mathbf{w}}$.

Instead, we use an iterative *gradient ascent* method to compute the optimal value.

The gradient ascent method relies on the gradient of the log-likelihood function, which can be obtained by taking its partial derivative with respect to $\tilde{\mathbf{w}}$, as follows:

$$\nabla(\tilde{\mathbf{w}}) = \frac{\partial}{\partial \tilde{\mathbf{w}}} \left\{ \ln(L(\tilde{\mathbf{w}})) \right\} = \frac{\partial}{\partial \tilde{\mathbf{w}}} \left\{ \sum_{i=1}^n y_i \cdot \ln(\theta(z_i)) + (1 - y_i) \cdot \ln(\theta(-z_i)) \right\}$$

where $z_i = \tilde{\mathbf{w}}^T \tilde{\mathbf{x}}_i$.

After solving the derivative:

$$\nabla(\tilde{\mathbf{w}}) = \sum_{i=1}^n (y_i - \theta(\tilde{\mathbf{w}}^T \tilde{\mathbf{x}}_i)) \cdot \tilde{\mathbf{x}}_i$$

Maximum Likelihood Estimation

The gradient ascent method starts at $\tilde{\mathbf{w}}^0$.

At each step t , the method moves in the direction of steepest ascent, which is given by the gradient vector.

Given the current $\tilde{\mathbf{w}}^t$, the next estimate is:

$$\tilde{\mathbf{w}}^{t+1} = \tilde{\mathbf{w}}^t + \eta \cdot \nabla(\tilde{\mathbf{w}}^t)$$

$\eta > 0$ is the *learning rate*. It should not be too large, otherwise the estimates will vary wildly from one iteration to the next, and it should not be too small, otherwise it will take a long time to converge.

At the optimal value of $\tilde{\mathbf{w}}$, the gradient will be zero, i.e., $\nabla(\tilde{\mathbf{w}}) = 0$, as desired.

Stochastic Gradient Ascent

The gradient ascent method computes the gradient by considering all the data points, and it is therefore called *batch* gradient ascent.

For large datasets, it is typically much faster to compute the gradient by considering only one (randomly chosen) point at a time, which is called *stochastic gradient ascent* (SGA).

$$\nabla(\tilde{\mathbf{w}}, \tilde{\mathbf{x}}_i) = (y_i - \theta(\tilde{\mathbf{w}}^T \tilde{\mathbf{x}}_i)) \cdot \tilde{\mathbf{x}}_i$$

Once the model has been trained, we can predict the response for any new augmented test point $\tilde{\mathbf{z}}$ as follows:

$$\hat{y} = \begin{cases} 1 & \text{if } \theta(\tilde{\mathbf{w}}^T \tilde{\mathbf{z}}) \geq 0.5 \\ 0 & \text{if } \theta(\tilde{\mathbf{w}}^T \tilde{\mathbf{z}}) < 0.5 \end{cases}$$

Logistic Regression: Stochastic Gradient Ascent

LogisticRegression-SGA (D, η, ϵ):

```
1 foreach  $\mathbf{x}_i \in D$  do  $\tilde{\mathbf{x}}_i^T \leftarrow (1 \quad \mathbf{x}_i^T)$  // map to  $\mathbb{R}^{d+1}$ 
2  $t \leftarrow 0$  // step/iteration counter
3  $\tilde{\mathbf{w}}^0 \leftarrow (0, \dots, 0)^T \in \mathbb{R}^{d+1}$  // initial weight vector
4 repeat
5    $\tilde{\mathbf{w}} \leftarrow \tilde{\mathbf{w}}^t$  // make a copy of  $\tilde{\mathbf{w}}^t$ 
6   foreach  $\tilde{\mathbf{x}}_i \in \tilde{D}$  in random order do
7      $\nabla(\tilde{\mathbf{w}}, \tilde{\mathbf{x}}_i) \leftarrow (y_i - \theta(\tilde{\mathbf{w}}^T \tilde{\mathbf{x}}_i)) \cdot \tilde{\mathbf{x}}_i$  // gradient at  $\tilde{\mathbf{x}}_i$ 
8      $\tilde{\mathbf{w}} \leftarrow \tilde{\mathbf{w}} + \eta \cdot \nabla(\tilde{\mathbf{w}}, \tilde{\mathbf{x}}_i)$  // update estimate for  $\tilde{\mathbf{w}}$ 
9    $\tilde{\mathbf{w}}^{t+1} \leftarrow \tilde{\mathbf{w}}$  // update  $\tilde{\mathbf{w}}^{t+1}$ 
10   $t \leftarrow t + 1$ 
11 until  $\|\tilde{\mathbf{w}}^t - \tilde{\mathbf{w}}^{t-1}\| \leq \epsilon$ 
```


Logistic Regression

Example

Figure shows the Iris principal components data.

X_1 and X_2 are the independent attributes and represent the first two principal components.

Y is the binary response variable and represents the type of Iris flower; $Y = 1$ corresponds to Iris-virginica, whereas $Y = 0$ corresponds to the two other Iris types, namely Iris-setosa and Iris-versicolor.



Linear Regression

Example


The plane of best fit in linear regression has the weight vector:

$$\tilde{\mathbf{w}} = (0.333, -0.167, 0.074)^T$$

$$\hat{y} = f(\tilde{\mathbf{x}}) = 0.333 - 0.167 \cdot x_1 + 0.074 \cdot x_2$$

Since Y is binary, we predict $y = 1$ if $f(\tilde{\mathbf{x}}) \geq 0.5$, and $y = 0$ otherwise.

Linear regression misclassifies 17 points, achieving 88.7% accuracy.



Logistic Regression


The fitted logistic model is given as

$$\tilde{\mathbf{w}} = (w_0, w_1, w_2)^T = (-6.79, -5.07, -3.29)^T$$

$$P(Y = 1|\tilde{\mathbf{x}}) = \theta(\tilde{\mathbf{w}}^T \tilde{\mathbf{x}}) = \frac{1}{1 + \exp\{6.79 + 5.07 \cdot x_1 + 3.29 \cdot x_2\}}$$

Given $\tilde{\mathbf{x}}$, if $P(Y = 1|\tilde{\mathbf{x}}) \geq 0.5$, then we predict $\hat{y} = 1$, otherwise we predict $\hat{y} = 0$.

Logistic regression misclassifies only 5 points, achieving 96.7% accuracy.



Multiclass Logistic Regression

We now generalize logistic regression to the case when Y can take on K distinct nominal categorical values called *classes*, i.e., $Y \in \{c_1, c_2, \dots, c_K\}$.

We model Y as a K -dimensional multivariate Bernoulli random variable. Since Y can assume only one of the K values, we use the *one-hot encoding* approach to map each categorical value c_i to the K -dimensional binary vector

$$\mathbf{e}_i = (\underbrace{0, \dots, 0}_{i-1}, 1, \underbrace{0, \dots, 0}_{K-i})^T$$

whose i th element $e_{ii} = 1$, and all other elements $e_{ij} = 0$, so that $\sum_{j=1}^K e_{ij} = 1$.

The probability mass function for $\mathbf{Y}$ given $\tilde{\mathbf{X}} = \tilde{\mathbf{x}}$ is

$$P(\mathbf{Y} = \mathbf{e}_i | \tilde{\mathbf{X}} = \tilde{\mathbf{x}}) = \pi_i(\tilde{\mathbf{x}}), \text{ for } i = 1, 2, \dots, K$$

There are K unknown parameters, which must satisfy the following constraint:

$$\sum_{i=1}^K \pi_i(\tilde{\mathbf{x}}) = \sum_{i=1}^K P(\mathbf{Y} = \mathbf{e}_i | \tilde{\mathbf{X}} = \tilde{\mathbf{x}}) = 1$$

Multiclass Logistic Regression

Given that only one element of $\mathbf{Y}$ is 1, the PMF of $\mathbf{Y}$ is:

$$P(\mathbf{Y}|\tilde{\mathbf{X}} = \tilde{\mathbf{x}}) = \prod_{j=1}^K (\pi_j(\tilde{\mathbf{x}}))^{\mathbf{Y}_j}$$

We select c_K as a reference or base class, and consider the log-odds ratio of the other classes w.r.t. c_K .

We assume these log-odd ratios are linear in $\tilde{\mathbf{X}}$, but $\tilde{\omega}_i$ is specific to for class c_i :

$$\begin{aligned}\ln(\text{odds}(\mathbf{Y} = \mathbf{e}_i|\tilde{\mathbf{X}} = \tilde{\mathbf{x}})) &= \ln\left(\frac{P(\mathbf{Y} = \mathbf{e}_i|\tilde{\mathbf{X}} = \tilde{\mathbf{x}})}{P(\mathbf{Y} = \mathbf{e}_K|\tilde{\mathbf{X}} = \tilde{\mathbf{x}})}\right) = \ln\left(\frac{\pi_i(\tilde{\mathbf{x}})}{\pi_K(\tilde{\mathbf{x}})}\right) = \tilde{\omega}_i^T \tilde{\mathbf{x}} \\ &= \omega_{i0} \cdot x_0 + \omega_{i1} \cdot x_1 + \cdots + \omega_{id} \cdot x_d\end{aligned}$$

where $\omega_{i0} = \beta_i$ is the true bias value for class c_i .

Multiclass Logistic Regression

Setting $\tilde{\omega}_K = 0$, we have $\exp\{\tilde{\omega}_K^T \tilde{\mathbf{x}}\} = 1$, and thus we can write the full model for multiclass logistic regression as follows:

$$\pi_i(\tilde{\mathbf{x}}) = \frac{\exp\{\tilde{\omega}_i^T \tilde{\mathbf{x}}\}}{\sum_{j=1}^K \exp\{\tilde{\omega}_j^T \tilde{\mathbf{x}}\}}, \quad \text{for all } i = 1, 2, \dots, K$$

This function is also called the *softmax* function.

When $K = 2$, this formulation yields exactly the same model as in binary logistic regression.

Note that the choice of the reference class is not important, since we can derive the log-odds ratio for any two classes c_i and c_j .

Maximum Likelihood Estimation

To find the K sets of regression weight vectors $\tilde{\mathbf{w}}_i$, for $i = 1, 2, \dots, K$, we use the gradient ascent approach to maximize the log-likelihood function. The likelihood of the data is given as

$$L(\tilde{\mathbf{W}}) = P(\mathbf{Y}|\tilde{\mathbf{W}}) = \prod_{i=1}^n P(\mathbf{y}_i|\tilde{\mathbf{X}} = \tilde{\mathbf{x}}_i) = \prod_{i=1}^n \prod_{j=1}^K (\pi_j(\tilde{\mathbf{x}}_i))^{y_{ij}}$$

where $\tilde{\mathbf{W}} = \{\tilde{\mathbf{w}}_1, \tilde{\mathbf{w}}_2, \dots, \tilde{\mathbf{w}}_K\}$ is the set of K weight vectors.

The log-likelihood is then given as:

$$\ln(L(\tilde{\mathbf{W}})) = \sum_{i=1}^n \sum_{j=1}^K y_{ij} \cdot \ln(\pi_j(\tilde{\mathbf{x}}_i)) = \sum_{i=1}^n \sum_{j=1}^K y_{ij} \cdot \ln \left(\frac{\exp\{\tilde{\mathbf{w}}_j^T \tilde{\mathbf{x}}_i\}}{\sum_{a=1}^K \exp\{\tilde{\mathbf{w}}_a^T \tilde{\mathbf{x}}_i\}} \right)$$

Note that the negative of the log-likelihood function can be regarded as an error function, commonly known as *cross-entropy error*.

For stochastic gradient ascent, we update the weight vectors by considering only one point at a time.

Maximum Likelihood Estimation

The gradient of the log-likelihood function w.r.t. $\tilde{\mathbf{w}}_j$ at a given point $\tilde{\mathbf{x}}_i$ is:

$$\nabla(\tilde{\mathbf{w}}_j, \tilde{\mathbf{x}}_i) = (y_{ij} - \pi_j(\tilde{\mathbf{x}}_i)) \cdot \tilde{\mathbf{x}}_i$$

which results in the following update rule for the j th weight vector:

$$\tilde{\mathbf{w}}_j^{t+1} = \tilde{\mathbf{w}}_j^t + \eta \cdot \nabla(\tilde{\mathbf{w}}_j^t, \tilde{\mathbf{x}}_i)$$

where $\tilde{\mathbf{w}}_j^t$ denotes the estimate of $\tilde{\mathbf{w}}_j$ at step t , and η is the learning rate.

Once the model has been trained, we can predict $\hat{y}$ for any new $\tilde{\mathbf{z}}$ as:

$$\hat{y} = \arg \max_{c_i} \{ \pi_i(\tilde{\mathbf{z}}) \} = \arg \max_{c_i} \left\{ \frac{\exp\{\tilde{\mathbf{w}}_i^T \tilde{\mathbf{z}}\}}{\sum_{j=1}^K \exp\{\tilde{\mathbf{w}}_j^T \tilde{\mathbf{z}}\}} \right\}$$

We evaluate the softmax function and the predicted $\hat{y}$ has the highest probability.

Multiclass Logistic Regression Algorithm

LogisticRegression-MultiClass (D, η, ϵ):

```
1 foreach  $(\mathbf{x}_i^T, y_i) \in D$  do
2    $\tilde{\mathbf{x}}_i^T \leftarrow (1 \quad \mathbf{x}_i^T)$  // map to  $\mathbb{R}^{d+1}$ 
3    $\mathbf{y}_i \leftarrow \mathbf{e}_j$  if  $y_i = c_j$  // map  $y_i$  to  $K$ -dim Bernoulli vector
4  $t \leftarrow 0$  // step/iteration counter
5 foreach  $j = 1, 2, \dots, K$  do  $\tilde{\mathbf{w}}_j^t \leftarrow (0, \dots, 0)^T \in \mathbb{R}^{d+1}$ 
6 repeat
7   foreach  $j = 1, 2, \dots, K-1$  do  $\tilde{\mathbf{w}}_j \leftarrow \tilde{\mathbf{w}}_j^t$  // copy  $\tilde{\mathbf{w}}_j^t$ 
8   foreach  $\tilde{\mathbf{x}}_i \in \tilde{D}$  in random order do
9     foreach  $j = 1, 2, \dots, K-1$  do
10       $\pi_j(\tilde{\mathbf{x}}_i) \leftarrow \exp\{\tilde{\mathbf{w}}_j^T \tilde{\mathbf{x}}_i\} / \sum_{a=1}^K \exp\{\tilde{\mathbf{w}}_a^T \tilde{\mathbf{x}}_i\}$ 
11       $\nabla(\tilde{\mathbf{w}}_j, \tilde{\mathbf{x}}_i) \leftarrow (y_{ij} - \pi_j(\tilde{\mathbf{x}}_i)) \cdot \tilde{\mathbf{x}}_i$  // gradient at  $\tilde{\mathbf{w}}_j$ 
12       $\tilde{\mathbf{w}}_j \leftarrow \tilde{\mathbf{w}}_j + \eta \cdot \nabla(\tilde{\mathbf{w}}_j, \tilde{\mathbf{x}}_i)$  // update estimate for  $\tilde{\mathbf{w}}_j$ 
13   foreach  $j = 1, 2, \dots, K-1$  do  $\tilde{\mathbf{w}}_j^{t+1} \leftarrow \tilde{\mathbf{w}}_j$  // update  $\tilde{\mathbf{w}}_j^{t+1}$ 
14    $t \leftarrow t + 1$ 
15 until  $\sum_{j=1}^{K-1} \|\tilde{\mathbf{w}}_j^t - \tilde{\mathbf{w}}_j^{t-1}\| \leq \epsilon$ 
```

Multiclass Logistic Regression


Example

Consider the 2D Iris PCA dataset, $n = 150$.

Y takes on three values: $Y = c_1$: Iris-setosa ($\square$), $Y = c_2$: Iris-versicolor ($\circ$) and $Y = c_3$: Iris-virginica ($\triangle$).

$Y = c_1$ to $\mathbf{e}_1 = (1, 0, 0)^T$, $Y = c_2$ to $\mathbf{e}_2 = (0, 1, 0)^T$ and $Y = c_3$ to $\mathbf{e}_3 = (0, 0, 1)^T$.

All the points actually lie in the (X_1, X_2) plane, but c_1 and c_2 are shown displaced along Y with respect to the base class c_3 purely for illustration purposes.



Multiclass Logistic Regression


Example

We use $Y = c_3$ as the reference or base class. The fitted model is:

$$\tilde{\mathbf{w}}_1 = (-3.52, 3.62, 2.61)^T \quad \tilde{\mathbf{w}}_2 = (-6.95, -5.18, -3.40)^T \quad \tilde{\mathbf{w}}_3 = (0, 0, 0)^T$$

The decision surface corresponding to c_1 is:

$$\pi_1(\tilde{\mathbf{x}}) = \frac{\exp\{\tilde{\mathbf{w}}_1^T \tilde{\mathbf{x}}\}}{1 + \exp\{\tilde{\mathbf{w}}_1^T \tilde{\mathbf{x}}\} + \exp\{\tilde{\mathbf{w}}_2^T \tilde{\mathbf{x}}\}}$$



Multiclass Logistic Regression


Example

We use $Y = c_3$ as the reference or base class. The fitted model is:

$$\tilde{\mathbf{w}}_1 = (-3.52, 3.62, 2.61)^T \quad \tilde{\mathbf{w}}_2 = (-6.95, -5.18, -3.40)^T \quad \tilde{\mathbf{w}}_3 = (0, 0, 0)^T$$

The decision surface corresponding to c_2 is:

$$\pi_2(\tilde{\mathbf{x}}) = \frac{\exp\{\tilde{\mathbf{w}}_2^T \tilde{\mathbf{x}}\}}{1 + \exp\{\tilde{\mathbf{w}}_1^T \tilde{\mathbf{x}}\} + \exp\{\tilde{\mathbf{w}}_2^T \tilde{\mathbf{x}}\}}$$



Multiclass Logistic Regression


Example

We use $Y = c_3$ as the reference or base class. The fitted model is:

$$\tilde{\mathbf{w}}_1 = (-3.52, 3.62, 2.61)^T \quad \tilde{\mathbf{w}}_2 = (-6.95, -5.18, -3.40)^T \quad \tilde{\mathbf{w}}_3 = (0, 0, 0)^T$$

The decision surface corresponding to c_3 is:

$$\pi_3(\tilde{\mathbf{x}}) = \frac{1}{1 + \exp\{\tilde{\mathbf{w}}_1^T \tilde{\mathbf{x}}\} + \exp\{\tilde{\mathbf{w}}_2^T \tilde{\mathbf{x}}\}}$$



Multiclass Logistic Regression

Example

The training set accuracy is 96.7%, since it misclassifies only five points (shown in dark gray).


For example, for the point $\tilde{\mathbf{x}} = (1, -0.52, -1.19)^T$, we have:

$$\pi_1(\tilde{\mathbf{x}}) = 0$$

$$\pi_2(\tilde{\mathbf{x}}) = 0.448$$

$$\pi_3(\tilde{\mathbf{x}}) = 0.552$$

$\hat{y} = \arg \max_{c_i} \{\pi_i(\tilde{\mathbf{x}})\} = c_3$, whereas the true class is $y = c_2$.



Data Mining and Machine Learning

A Master Course

Master program on Data Science

Tran Trong Hieu¹

¹VNU - University of Engineering and Technology

Chapter 24: Logistic Regression